


CS122A: Introduction to Data Management

Lecture #5 (E-R \rightarrow Relational, Cont.)

Instructor: Chen Li

Logical DB Design: ER to Relational (Review)

❖ Entity sets to tables:




```
CREATE TABLE Employees  
  (ssn CHAR(11),  
   name CHAR(20),  
   lot INTEGER,  
   PRIMARY KEY (ssn))
```

Relationship Sets to Tables (Review)



- ❖ In translating a relationship set to a relation, attributes of the relation must include:
 - Keys for each participating entity set (as foreign keys).
 - Note: This set of attributes forms a *superkey* for the relation.
 - All descriptive attributes.

```
CREATE TABLE Works_In(  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (ssn, did),  
    FOREIGN KEY (ssn)  
        REFERENCES Employees,  
    FOREIGN KEY (did)  
        REFERENCES Departments)
```




Key Constraints (Review)


- ❖ Each dept has at most one manager, according to the key constraint on Manages.




1-to-1



1-to Many



Many-to-1



Many-to-Many

Translation to relational model?

Translating ER Diagrams with Key Constraints

❖ Map the relationship to a table (Manages):

- Note that **did** is the key now!
- Separate tables for Employees and Departments.

```
CREATE TABLE Manages (  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    FOREIGN KEY (did) REFERENCES Departments)
```

vs.


❖ But, since *each* department has a *unique* manager, we could choose to fold Manages right into Departments.

(Q: Why do that...?)

```
CREATE TABLE Departments2 (  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    mgr_ssn CHAR(11),  
    mgr_since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (mgr_ssn) REFERENCES Employees)
```

***Note:** The relationship info has been pushed to the N-side's entity table!*

Properly Reflecting Key Constraints




❖ So what are the translated relationship table keys (*etc.*) when...

- FooBar is M:N? → FooBar(fooId, barId, baz)
- FooBar is N:1? → FooBar(fooId, barId, baz)
- FooBar is 1:N? → FooBar(fooId, barId, baz)
- FooBar is 1:1? → FooBar(fooId, barId, baz) (Note: unique)

Review: Participation Constraints

- ❖ Does every department have a manager?
 - If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).
 - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!!)




Participation Constraints in SQL

- ❖ We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to the use of *triggers*).

```
CREATE TABLE Department2 (  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    mgr_ssn CHAR(11) NOT NULL,  
    mgr_since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (mgr_ssn) REFERENCES Employees,  
    ON DELETE NO ACTION)
```


Review: Weak Entities

- ❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.



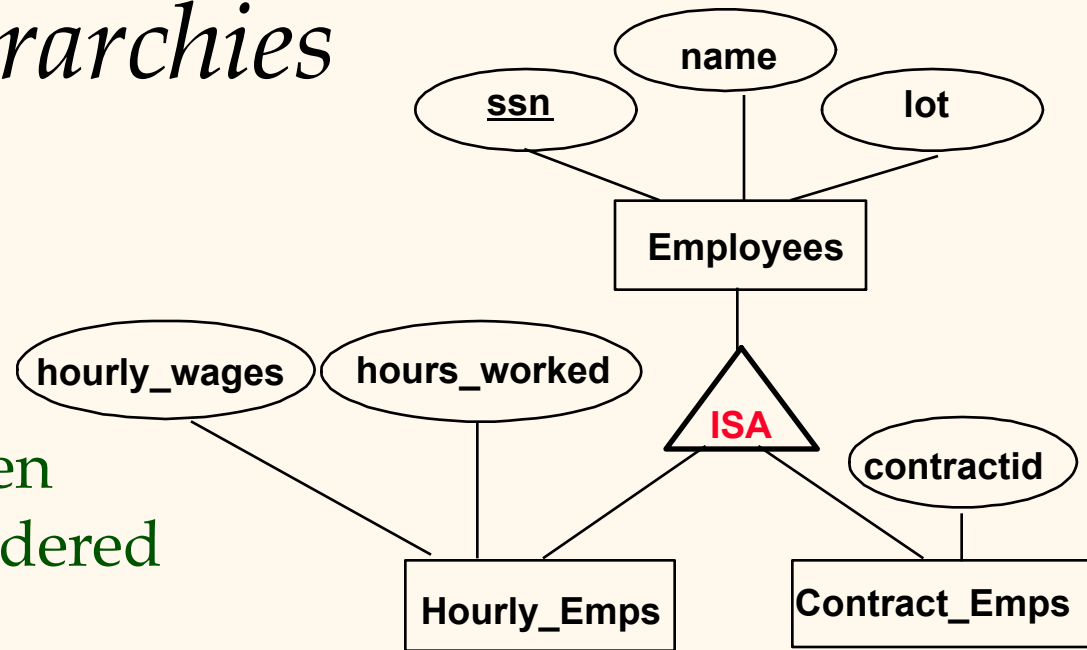
Translating Weak Entity Sets

- ❖ Weak entity set and identifying relationship set are translated into a single table.
 - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dependents2 (  
    pname CHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (pname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

Review: ISA Hierarchies

- ❖ As in C++, or other PLs, attributes are inherited.
- ❖ If we declare A **ISA** B, then every A entity is also considered to be a B entity.



- ❖ *Overlap constraints*: Can Joe be an Hourly_Emps as well as a Contract_Emps entity? (*Allowed/disallowed*)
- ❖ *Covering constraints*: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? (*Yes/no*)

Translating ISA Hierarchies to Relations

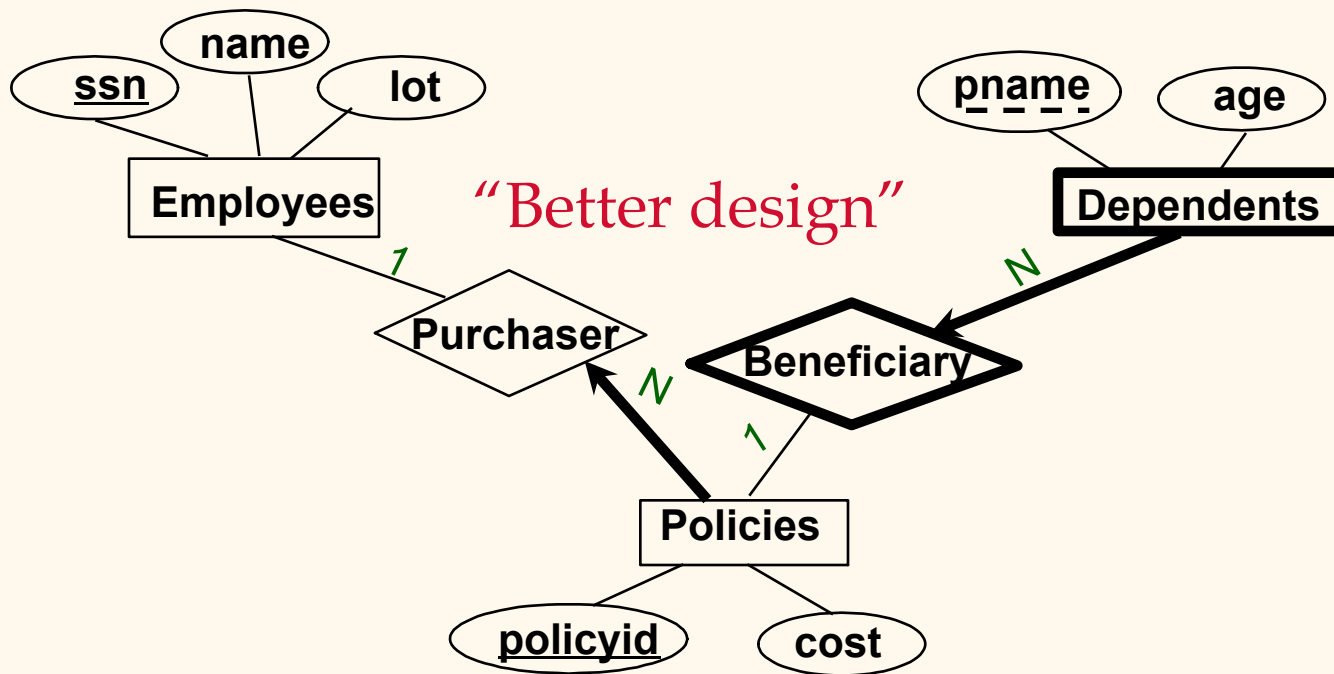
❖ *Most general approach:*

- 3 relations: Employees, Hourly_Emps and Contract_Emps.
 - *Hourly_Emps*: Every employee recorded in Employees. For hourly emps, *extra* info recorded in Hourly_Emps (*hourly_wages*, *hours_worked*, *ssn*); delete Hourly_Emps tuple if referenced Employees tuple is deleted.
 - Queries about all employees easy; those involving just Hourly_Emps require a join to get the extra attributes.

❖ *An alternative: Hourly_Emps and Contract_Emps.*

- *Hourly_Emps*: *ssn*, *name*, *lot*, *hourly_wages*, *hours_worked*.
- Each employee must be in one of these two subclasses.
(Q: Can we always do this, then? A: Not w/o redundancy!)

Review: Binary vs. Ternary Relationships



Binary vs. Ternary Relationships (Contd.)

- ❖ The key constraints let us combine *Purchaser with Policies* and *Beneficiary with Dependents*.


```
CREATE TABLE Policies (  
  policyid INTEGER,  
  cost REAL,  
  emp_ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (policyid).  
  FOREIGN KEY (emp_ssn) REFERENCES Employees  
    ON DELETE CASCADE)
```

- ❖ Participation constraints lead to **NOT NULL** constraints.
(Note: Primary key attributes are NOT NULL as well – check documentation to see if that's implicit or explicit!)

```
CREATE TABLE Dependents (  
  pname CHAR(20),  
  age INTEGER,  
  policyid INTEGER,  
  PRIMARY KEY (pname, policyid),  
  FOREIGN KEY (policyid) REFERENCES Policies  
    ON DELETE CASCADE)
```

Review: Binary vs. Ternary Relationships

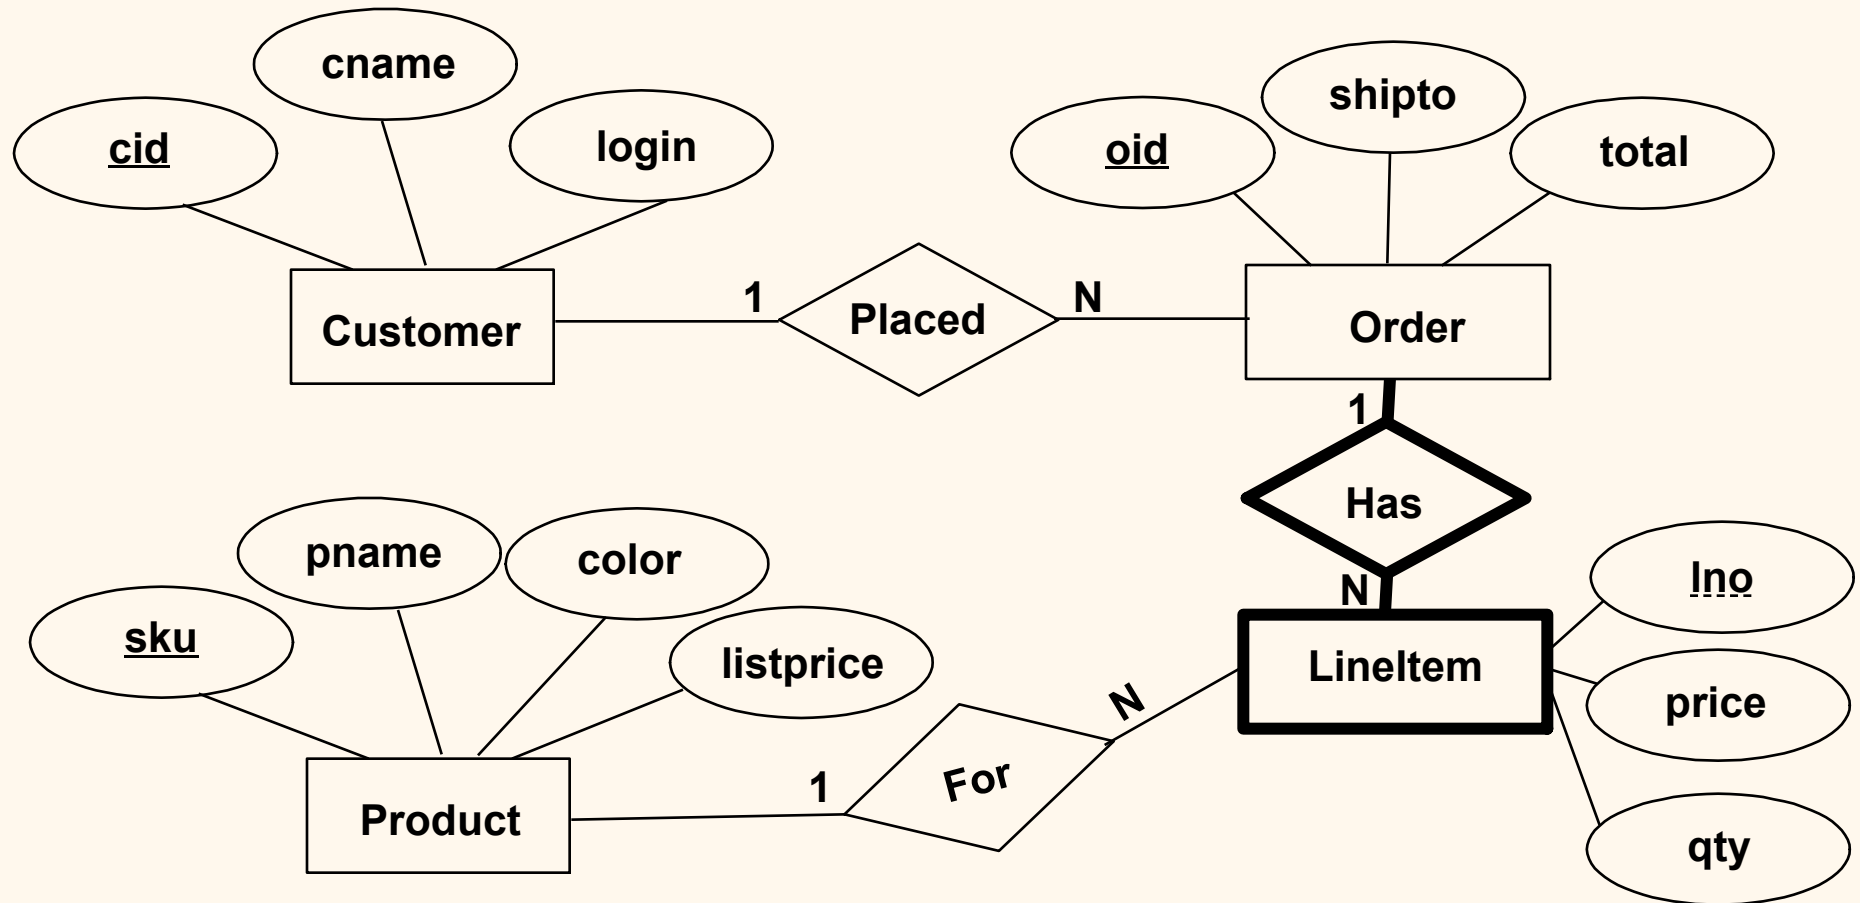
```
CREATE TABLE Employees (  
  ssn CHAR(11),  
  name CHAR(20),  
  lot INTEGER,  
  PRIMARY KEY (ssn))
```



```
CREATE TABLE Policies (  
  policyid INTEGER,  
  cost REAL,  
  emp_ssn CHAR(11) NOT NULL,  
  PRIMARY KEY (policyid).  
  FOREIGN KEY (emp_ssn)  
    REFERENCES Employees  
    ON DELETE CASCADE)
```

```
CREATE TABLE Dependents (  
  pname CHAR(20),  
  age INTEGER,  
  policyid INTEGER,  
  PRIMARY KEY (pname, policyid),  
  FOREIGN KEY (policyid)  
    REFERENCES Policies  
    ON DELETE CASCADE)
```

An Example: Putting It Together



Putting It Together (Cont'd.)

```
CREATE TABLE Customer (  
  cid INTEGER,  
  cname VARCHAR(50),  
  login VARCHAR(20)  
    NOT NULL,  
  PRIMARY KEY (cid),  
  UNIQUE (login))
```

```
CREATE TABLE Product (  
  sku INTEGER,  
  pname VARCHAR(100),  
  color VARCHAR(20),  
  listprice DECIMAL(8,2),  
  PRIMARY KEY (sku))
```

```
CREATE TABLE Order (  
  oid INTEGER,  
  custid INTEGER,  
  shipto VARCHAR(200),  
  total DECIMAL(8,2),  
  PRIMARY KEY (oid),  
  FOREIGN KEY (custid) REFERENCES Customer))
```

```
CREATE TABLE LineItem (  
  oid INTEGER,  
  lno INTEGER,  
  price DECIMAL(8,2),  
  qty INTEGER,  
  sku INTEGER,  
  PRIMARY KEY (oid, lno),  
  FOREIGN KEY (oid) REFERENCES Order  
    ON DELETE CASCADE),  
  FOREIGN KEY (sku) REFERENCES Product))
```

Putting It Together (Cont'd.)

Customer

cid	cname	login
1	Smith, James	jsmith@aol.com
2	White, Susan	<u>suzie@gmail.com</u>
3	Smith, James	js@hotmail.com

Product

sku	pname	color	listprice
123	Frozen DVD	null	24.95
456	Graco Twin Stroller	green	199.99
789	Moen Kitchen Sink	black	350.00

Order

oid	custid	shipto	total
1	3	J. Smith, 1 Main St., USA	199.95
2	1	Mrs. Smith, 3 State St., USA	300.00

LineItem

oid	lno	price	qty	item
1	1	169.95	1	456
1	2	15.00	2	123
2	1	300.00	1	789

SQL Views

- ❖ A view is just a relation, but we store its *definition* rather than storing the (materialized) set of tuples.

```
CREATE VIEW YoungActiveStudents (name, grade)
AS SELECT S.name, E.grade
FROM Students S, Enrolled E
WHERE S.sid = E.sid and S.age < 21
```

- ❖ Views can be dropped using the **DROP VIEW** command.
 - How to handle **DROP TABLE** if there's a view on the table?
 - DROP TABLE command has options to let the user specify this.

Views and Security

- ❖ Views can be used to present necessary information (or a summary) while hiding some details in underlying relation(s).
 - Given YoungStudents, but not Students or Enrolled, we can find students *S* who have are enrolled, but not the *cid*'s of the courses they are enrolled in.
- ❖ Other view uses in our ER translation context might include:
 - Derived attributes, e.g., age (vs. birthdate)
 - Simplifying/eliminating join paths (for SQL)

Relational Model and E-R Schema Translation: Summary

- ❖ A tabular representation of data.
- ❖ Simple and intuitive, also widely used.
- ❖ Integrity constraints can be specified by the DBA based on application semantics. DBMS then checks for violations.
 - Two important ICs: Primary and foreign keys (PKs, FKs).
 - In addition, we *always* have domain constraints.
- ❖ Powerful and natural query languages exist (soon!)
- ❖ Rules to translate E-R to relational model
 - Can be done by a human, or automatically (using a tool)