

CS122A: Introduction to Data Management


Lecture #12: Relational DB Design Theory (1)

Instructor: Chen Li

Given a Relational Schema...

- ❖ How do I know if my relational schema is a “good” logical database design or not?
 - What might make it “not good”?
 - How can I fix it, if indeed it’s “not good”?
 - How “good” is it, after I’ve fixed it?
- ❖ Note that your relational schema might have come from one of several places
 - You started from an E-R model (but maybe that model was “wrong” in some way?)
 - You went straight to relational in the first place
 - It’s not your schema – you inherited it! 😊

Ex: Wisconsin Sailing Club




Proposed schema design #1:

sid	sname	rating	age	date	bid	bname	color
22	Dustin	7	45.0	10/10/98	101	Interlake	blue
22	Dustin	7	45.0	10/10/98	102	Interlake	red
22	Dustin	7	45.0	10/8/98	103	Clipper	green
22	Dustin	7	45.0	10/7/98	104	Marine	red
31	Lubber	8	55.5	11/10/98	102	Interlake	red
31	Lubber	8	55.5	11/6/98	103	Clipper	green
31	Lubber	8	55.5	11/12/98	104	Marine	red
...

Q: Do you think this is a “good” design? (Why or why not?)

Ex: Wisconsin Sailing Club



Proposed schema design #2:

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
...


sid	bid	date
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
...

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Q: What about *this* design?

- Is #2 “better than #1...?”
Explain!
- Is it a “best” design?
- How can we go from design #1 to this one?

Ex: Wisconsin Sailing Club



Proposed schema design #3:

sid	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
...

sid	bid	date
22	101	10/10/98
22	102	10/10/98
22	103	10/8/98
22	104	10/7/98
31	102	11/10/98
31	103	11/6/98
31	104	11/12/98
...

bid	bname
101	Interlake
102	Interlake
103	Clipper
104	Marine

bid	color
101	blue
102	red
103	green
104	red

Q: What about *this* design?

- Is #3 “better” or “worse” than #2...?
- What sort of tradeoffs do you see between the two?

The Evils of Redundancy

- ❖ *Redundancy* is at the root of several problems associated with relational schemas:
 - Redundant storage
 - Insert/delete/update anomalies
- ❖ *Functional dependencies* can help in identifying problem schemas and suggesting refinements.
- ❖ Main refinement technique: *decomposition*, e.g., replace $R(ABCD)$ with $R1(AB) + R2(BCD)$.
- ❖ Decomposition should be used judiciously:
 - Is there reason to decompose a relation?
 - Does the decomposition cause any problems?

Good rule to follow:
“One fact, one place!”

Functional Dependencies (FDs)

- ❖ A functional dependency $X \rightarrow Y$ holds over relation R if, for every allowable instance r of R:
 - For $t1$ and $t2$ in r , $t1.X = t2.X$ implies $t1.Y = t2.Y$
 - I.e., given two tuples in r , if the X values agree, then the Y values must also agree. (X and Y can be *sets* of attributes.)
- ❖ An FD is a statement about *all* allowable relations.
 - Identified based on application semantics (similar to E-R).
 - Given some instance $r1$ of R, we can check to see if it violates some FD f , but we cannot tell if f holds over R!
- ❖ Saying K is a candidate key for R means that $K \rightarrow R$
 - Note: $K \rightarrow R$ does not require K to be *minimal*! If K minimal, then it is a candidate key.

Example: Constraints on an Entity Set

- ❖ Suppose you're given a relation called HourlyEmps:
 - HourlyEmps (ssn, name, lot, rating, hrly_wages, hrs_worked)
- ❖ Notation: We will denote this relation schema by simply listing the attributes: SNLRWH
 - This is really the *set* of attributes {S,N,L,R,W,H}.
 - Sometimes, we will refer to all attributes of a relation by using the relation name (e.g., HourlyEmps for SNLRWH).
- ❖ Suppose we also have some FDs on HourlyEmps:
 - *ssn* is the key: $S \rightarrow \text{SNLRWH}$
 - *rating* determines *hrly_wages*: $R \rightarrow W$

Example (Cont'd.)

Wages

R	W
8	10
5	7

HourlyEmps2

S	N	L	R	H
123-22-3666	Attishoo	48	8	40
231-31-5368	Smiley	22	8	30
131-24-3650	Smethurst	35	5	30
434-26-3751	Guldu	35	5	32
612-67-4134	Madayan	35	8	40

S	N	L	R	W	H
123-22-3666	Attishoo	48	8	10	40
231-31-5368	Smiley	22	8	10	30
131-24-3650	Smethurst	35	5	7	30
434-26-3751	Guldu	35	5	7	32
612-67-4134	Madayan	35	8	10	40

❖ Problems due to $R \rightarrow W$:

- Update anomaly: Can we change W in just the 1st tuple of SNLRWH?
- Insertion anomaly: What if we want to insert an employee and don't know the hourly wage for his rating?
- Deletion anomaly: If we delete all employees with rating 5, we lose the information about the wage for rating 5!

How about two smaller tables?

Reasoning About FDs

- ❖ Given some FDs, we can usually infer additional FDs:
 - $ssn \rightarrow did, did \rightarrow lot$ implies $ssn \rightarrow lot$
- ❖ An FD f is implied by a set of FDs F if f holds whenever all FDs in F hold.
 - $F^+ =$ closure of F is the set of all FDs that are implied by F .
- ❖ Armstrong's Axioms (X, Y, Z are sets of attributes):
 - Reflexivity: If $X \subseteq Y$, then $Y \rightarrow X$
 - Augmentation: If $X \rightarrow Y$, then $XZ \rightarrow YZ$ for any Z
 - Transitivity: If $X \rightarrow Y$ and $Y \rightarrow Z$, then $X \rightarrow Z$
- ❖ These are sound and complete inference rules for FDs!

Reasoning About FDs (Cont'd.)

(Recall: “two matching X’s always have the same Y”)

- ❖ A few additional rules (which follow from AA):
 - *Union*: If $X \rightarrow Y$ and $X \rightarrow Z$, then $X \rightarrow YZ$
 - *Decomposition*: If $X \rightarrow YZ$, then $X \rightarrow Y$ and $X \rightarrow Z$
- ❖ Example: **Contracts**(*cid,sid,jid,did,pid,qty,value*), and:
 - C is the key: $C \rightarrow CSJDPQV$
 - Project purchases each part using single contract: $JP \rightarrow C$
 - Dept purchases at most one part from a supplier: $SD \rightarrow P$
- ❖ $JP \rightarrow C, C \rightarrow CSJDPQV$ imply $JP \rightarrow CSJDPQV$
- ❖ $SD \rightarrow P$ implies $SDJ \rightarrow JP$
- ❖ $SDJ \rightarrow JP, JP \rightarrow CSJDPQV$ imply $SDJ \rightarrow CSJDPQV$

Reasoning About FDs (Examples)

Let's consider $R(ABCDE)$, $F = \{A \rightarrow B, B \rightarrow C, CD \rightarrow E\}$

❖ Let's work our way towards inferring F^+ ...

- | | | | |
|----------------------------|------------------------|------------------------|---------------------------------|
| (a) $A \rightarrow B$ | (b) $B \rightarrow C$ | (c) $CD \rightarrow E$ | (given) |
| (d) $A \rightarrow C$ | | | (a, b, and transitivity) |
| (e) $BD \rightarrow CD$ | | | (b and augmentation) |
| (f) $BD \rightarrow E$ | | | (e and transitivity) |
| (g) $AD \rightarrow CD$ | | | (d and augmentation) |
| (h) $AD \rightarrow E$ | | | (g and transitivity) |
| (i) $AD \rightarrow C$ | (j) $AD \rightarrow D$ | | (g and decomposition) |
| (j) $AD \rightarrow BD$ | | | (a and augmentation) |
| (l) $AD \rightarrow B$ | | | (k and decomposition) |
| (m) $AD \rightarrow A$ | | | (a and reflexivity) |
| (n) $AD \rightarrow ABCDE$ | | | (h, i, j, l, m, and union) |
- Candidate key!*
- If an attribute X is not on the RHS of *any* initial FD, X must be part of the key!

Reasoning About FDs (Cont'd.)

- ❖ Computing the closure of a set of FDs can be expensive. (Size of closure is exponential in #attrs!)
- ❖ Typically, we just want to check if a given FD $X \rightarrow Y$ is in the closure of a set of FDs F . An efficient check:
 - Compute attribute closure of X (denoted X^+) wrt F :
 - Set of all attributes A such that $X \rightarrow A$ is in F^+
 - There is a linear time algorithm to compute this.
 - Then check if Y is in X^+
- ❖ Does $F = \{A \rightarrow B, B \rightarrow C, C D \rightarrow E\}$ imply $A \rightarrow E$?
 - I.e.: is $A \rightarrow E$ in the closure F^+ ? Equivalently: Is E in A^+ ?


FDs & Redundancy

- ❖ Role of FDs in detecting redundancy:
 - Consider a relation R with 3 attributes, ABC.
 - **If no non-trivial FDs hold:** There is no redundancy here.
 - **Given $A \rightarrow B$:** Several tuples could have the same A value – and if so, then they'll all have the same B value as well!

Normal Forms

- ❖ Returning to the issue of schema refinement, the first question to ask is whether any refinement is needed!
- ❖ We will define various *normal form* (BCNF, 3NF etc.) based on the nature of FDs that hold
- ❖ Depending upon the normal form a relation is in, it has different level of redundancy
 - E.g., a BCNF relation has NO redundancy – will be clear soon!
- ❖ Checking for which normal form a relation is in will help us decide whether to decompose the relation
 - E.g., no point decomposing a BCNF relation!

→ Normal Forms ←




Some Terms and Definitions

- ❖ If X is part of a candidate key, we will say that X is a *prime attribute*.
- ❖ If X is not part of any candidate key, we will say that X is a *non-prime attribute*.
- ❖ If X (an attribute set) contains a candidate key, we will say that X is a *superkey*.
- ❖ $X \rightarrow Y$ can be pronounced as “ X determines Y ”, or “ Y is *functionally dependent* on X ”.
- ❖ $X \rightarrow Y$ is *trivial* if $Y \subseteq X$.

First Normal Form (1NF)

- ❖ Rel'n R is in **1NF** if all of its attributes are **atomic**.
 - No set-valued attributes! (1NF = “flat” 😊)
 - Usually goes *w/o* saying for relational model (but not for NoSQL systems, as we’ll see at the end of the quarter 😊).
 - *Ex:*




Second Normal Form (2NF)

Relation R is in **2NF** if

- It is in 1NF; and
- Each non-prime attribute is dependent on the *whole* of every candidate key

Violation:



Second Normal Form (2NF)

❖ *Ex:* Supplies(sno, sname, saddr, pno, pname, pcolor)

where: sno \rightarrow sname, sno \rightarrow saddr, pno \rightarrow pname, pno \rightarrow pcolor

Q1: What are the candidate keys for Supplies?

Q2: What are the prime attributes for Supplies?

Q3: Why is Supplies not in 2NF?

Q4: What's the fix?

Supplier(sno, sname, saddr)

Part(pno, pname, pcolor)

Supply(sno, pno)

A1: (sno, pno)

A2: sno, pno

A3: Each of its four
FDs violates 2NF!

Must not forget this!!
(Else “lossy join” !!)