Big Data / Web Warehousing

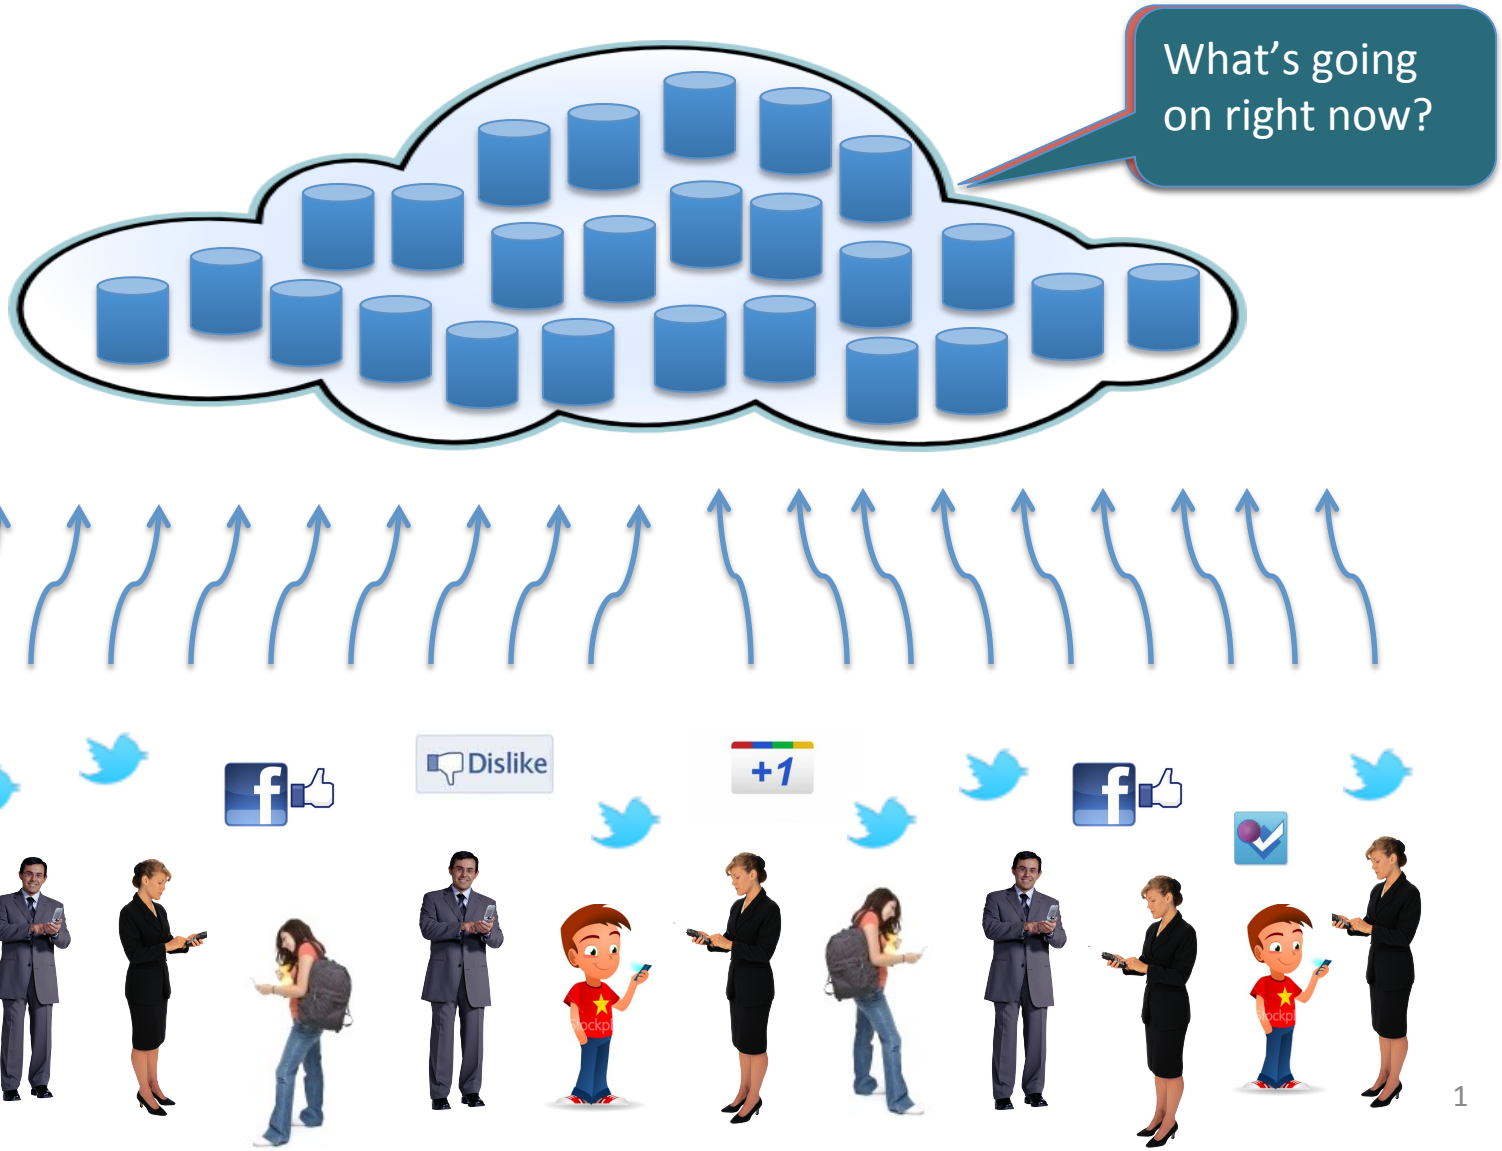

"Big Data" History

- Late 1990's brought a need to index and query the rapidly exploding content of the Web
  - DB technology tried but failed (*e.g.,* Inktomi)
  - Google, Yahoo! *et al* needed to do something
- Google responded by laying a new foundation
  - Google File System (GFS)
    - OS-level byte stream files spanning 1000's of machines
    - Three-way replication for fault-tolerance (availability)
  - MapReduce (MR) programming model
    - User functions: Map and Reduce (and optionally Combine)
    - *"Parallel programming for dummies"* – MR runtime does the heavy lifting via partitioned parallelism

(MapReduce: *Word Count* Example)

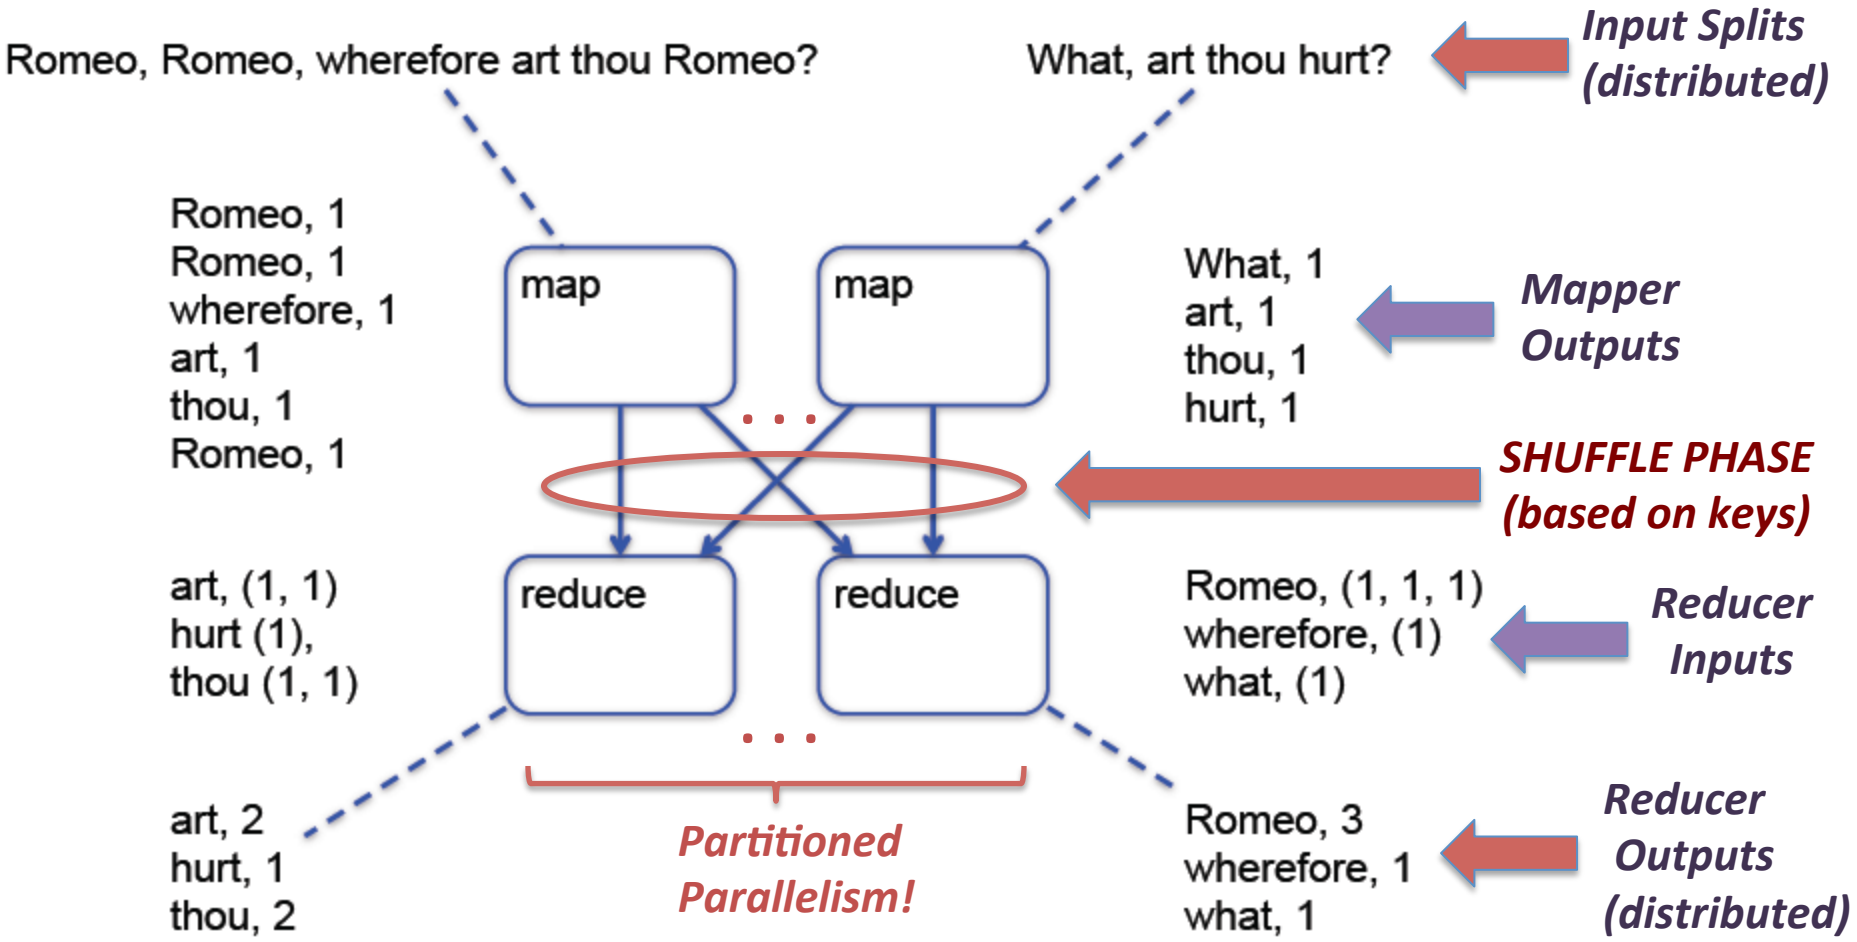

Today's "Big Data" Tangle

mongoDB

ly scalable, eventually consistent, distributed, structured key-value store.

AsterixDB: "One Size Fits a Bunch"

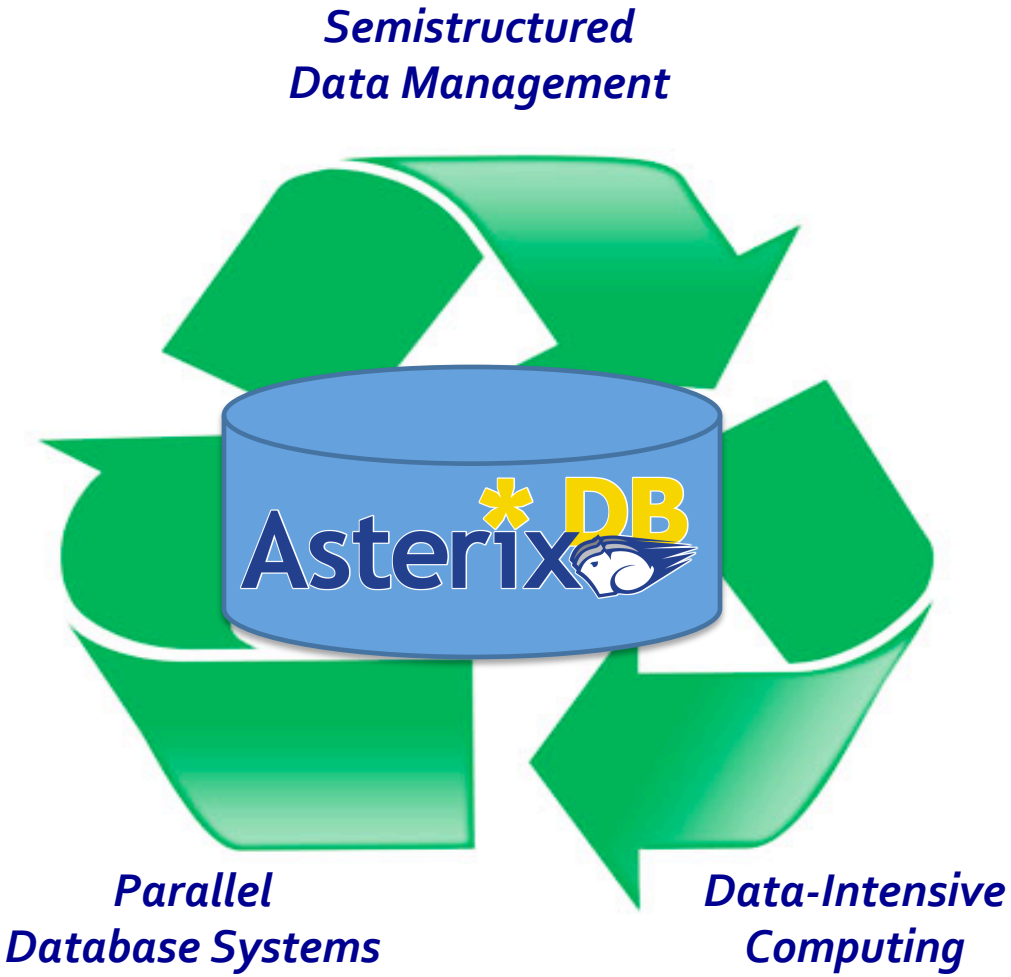

BDMS Desiderata:

- Flexible data model
- Efficient runtime
- Full query capability
- Cost proportional to task at hand (!)
- Designed for continuous data ingestion
- Support today's "Big Data data types"

Project Goals

- Build a new Big Data Management System (BDMS)
  - Run on large commodity clusters
  - Handle mass quantities of semistructured data
  - Openly *layered*, for selective reuse by others
  - Share with the community via *open source*
- Conduct scalable information systems research, e.g.,
  - Large-scale query processing and workload management
  - Highly scalable storage and index management
  - Fuzzy matching, spatial data, date/time data (all in parallel)
  - Novel support for "fast data" (both in and out)
- Train next generation of "Big Data" graduates

ASTERIX Data Model (ADM)

```
create dataverse TinySocial;
use dataverse TinySocial;

create type MugshotUserType as {
   id: int32,
   alias: string,
   name: string,
   user-since: datetime,
   address: {
      street: string,
      city: string,
      state: string,
      zip: string,
      country: string
   },
   friend-ids: {{ int32 }},
   employment: [EmploymentType]
}
```

```
create type EmploymentType as open {
   organization-name: string,
   start-date: date,
   end-date: date?
}

create dataset MugshotUsers(MugshotUserType)
                    primary key id;
```

Highlights include:

- JSON++ based data model
- Rich type support (spatial, temporal, …)
- Records, lists, bags
- ***Open vs. closed types***

ASTERIX Query Language (AQL)

- *Ex:* List the user name and messages sent by those users who joined the Mugshot social network in a certain time window:

```
for $user in dataset MugshotUsers
where $user.user-since >= datetime('2010-07-22T00:00:00')
  and $user.user-since <= datetime('2012-07-29T23:59:59')
return {
  "uname" : $user.name,
  "messages" :
      for $message in dataset MugshotMessages
      where $message.author-id = $user.id
      return $message.message
};
```

AsterixDB System Overview

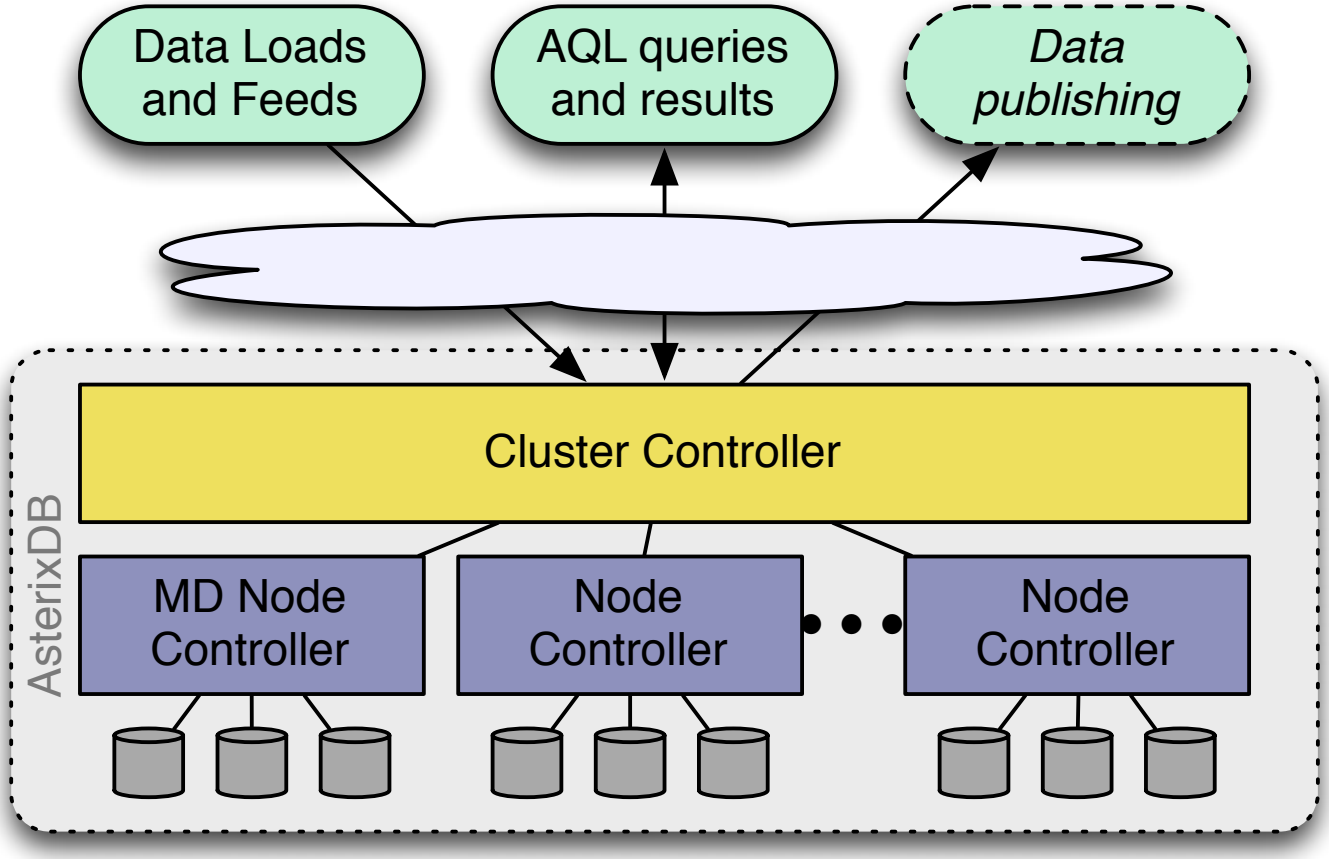

A prototype: interactive analytics and visualization of large data sets

http://cloudberry.ics.uci.edu/

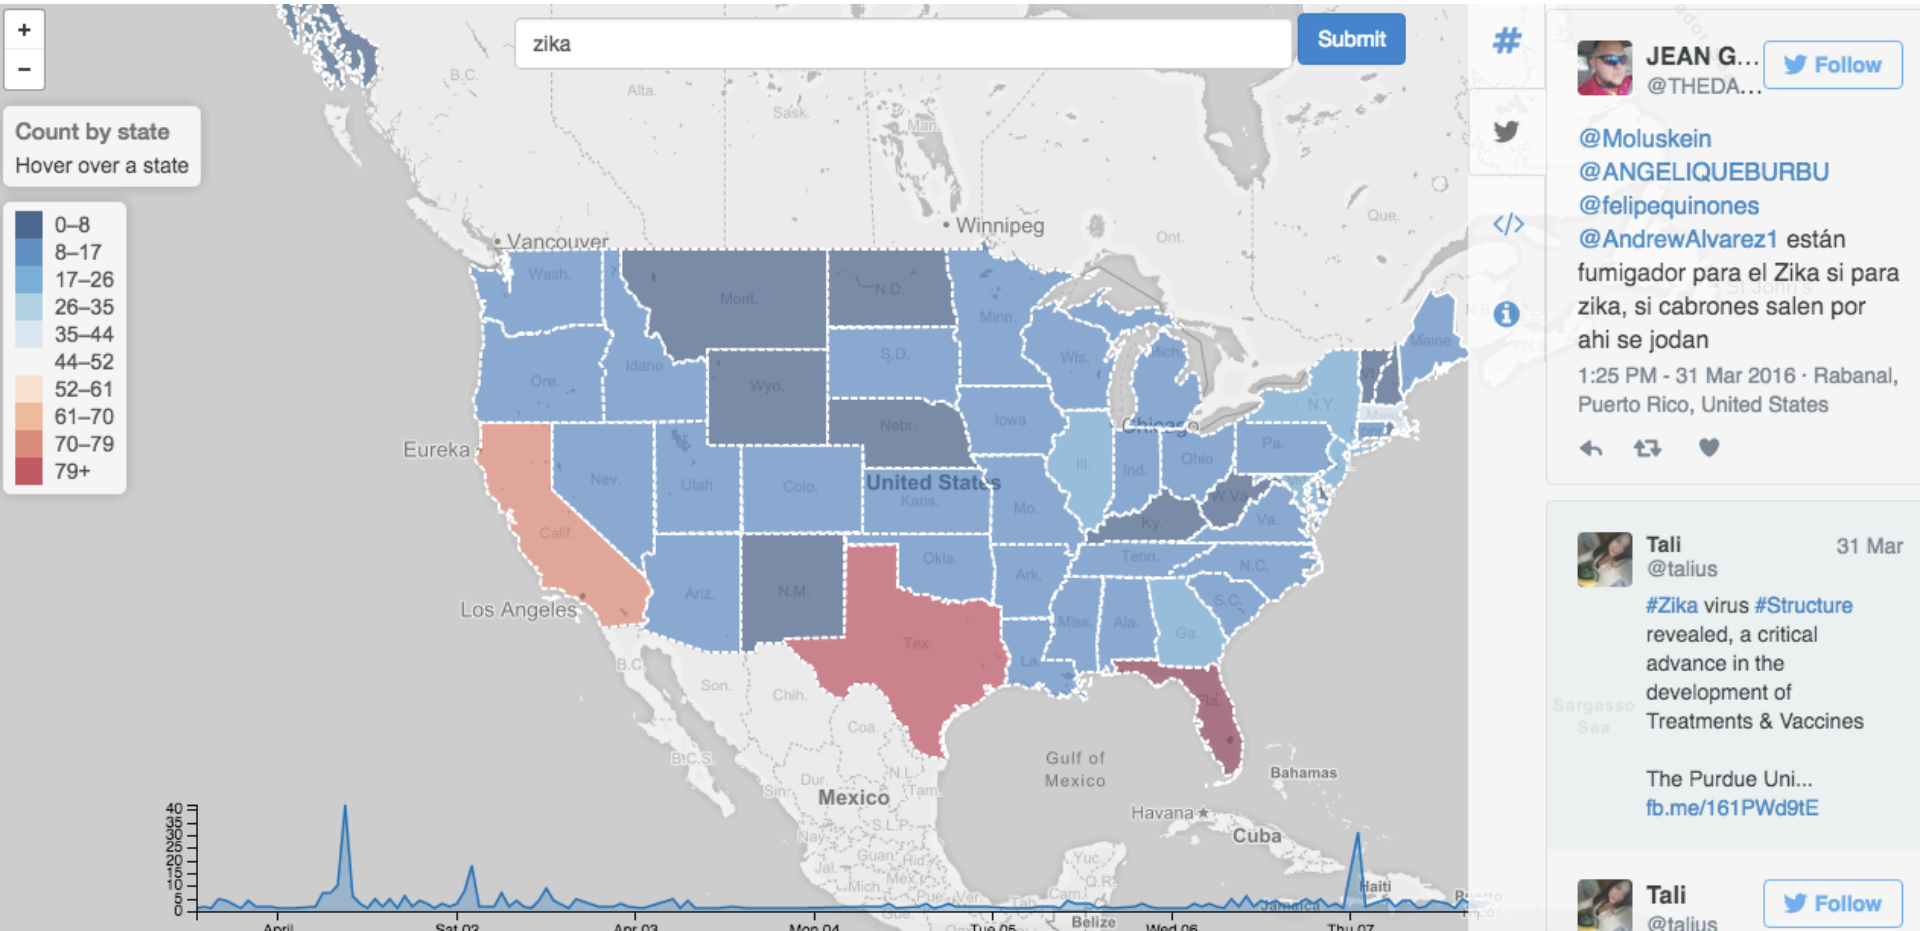

Apache AsterixDB project page:

https://asterixdb.apache.org/